

CSCI235 – Database System

Database Triggers

Sionggo Japit

sjapit@uow.edu.au

29 March 2026

Outline

What is a Database trigger?

Active database system

CREATE OR REPLACE TRIGGER statement

Statement database triggers

Row database triggers

Sample application

Problems with database triggers



What is a Database Trigger?

Database trigger is a piece of code stored in a data dictionary and automatically processed whenever a **pre-defined event** happens and **pre-defined condition** is satisfied

For example, we would like to automatically increase job level for all employees whose salary is above 100000

Modification of derived data

```
ON UPDATE OF EMPLOYEE.salary
  IF NEW.salary > 100000 THEN
    EXECUTE IncreaseJobLevel(NEW.enunder, NEW.salary);
  END IF;
```

For example, we would like to implement a data security rule saying that a salary cannot be updated over a weekend

Checking a constraint

```
ON UPDATE OF EMPLOYEE.salary
  IF EXTRACT(CURRENT_DATE, 'Day') IN ('Saturday', 'Sunday') THEN
    RAISE EXCEPTION 'Salary cannot be updated over a weekend !';
  END IF;
```

What is a Database Trigger?

For example, we would like to enforce a consistency constraint saying that a department cannot have more than 100 employees

```
ON INSERT INTO EMPLOYEE
```

Testing an event

```
SELECT COUNT(*)  
INTO total_employees  
FROM EMPLOYEE  
WHERE dname = NEW.dname;
```

Checking a consistency constraint

```
IF total_employees = 100 THEN  
    RAISE EXCEPTION 'Too many employees in %', NEW.dname);  
END IF;
```

Performing an action

In the example above we assume that a trigger **fires** and it is processed **before** `INSERT` statement

Sometimes it is more convenient to **fire** a trigger that verifies a consistency constraint **after** modification of a relational table and **before** `COMMIT` statement

This is why we have two temporal options for triggers: `BEFORE` and `AFTER`

What is a Database Trigger?


What do we need database triggers for ?

- To verify the consistency constraints
- To enforce the sophisticated database access controls
- To implement transparent event logging
- To generate the values of derived attributes
- To maintain replicated data in a distributed database
- To update the relational views

Active Database Systems provide functionalities for implementation of database triggers

Database Trigger

Active Database System



Active Database System

Active database system is a database system which is able to detect the **events** that have happened in a certain period of time and in the response to these **events** it is able to execute the **actions** when the **pre-defined conditions** are met

A logic of active database system is implemented as a collection of **Event-Condition-Action (ECA)** rules

In SQL **ECA** rule can be created with **CREATE TRIGGER** statement and it can be deleted with **DROP TRIGGER** statement

Syntax of **ECA rule**:

- (**EVENT**, **CONDITION**, **ACTION**)

Semantics of **ECA rule**:

- Whenever an **EVENT** happens and a **CONDITION** is satisfied then a database system performs an **ACTION**

Active Database System

A sample **event**

```
ON UPDATE OF EMPLOYEE.salary
```

Event

A sample **condition**

```
IF NEW.salary > 100000
```

Condition

A sample **Action**

```
EXECUTE IncreaseJobLevel(NEW.enumber, NEW.salary);
```

Trigger

CREATE OR REPLACE TRIGGER statement implements **ECA** rule



Database Trigger

Create or replace
Trigger statement

Create or Replace Trigger Statement

A sample `CREATE OR REPLACE TRIGGER` statement

```
CREATE OR REPLACE TRIGGER CheckBudget
```

Trigger name

Temporal option

```
BEFORE
```

Temporal option specification

Event

```
UPDATE OF budget ON DEPARTMENT
```

Event specification

Type of trigger, either `statement` or `row` trigger

```
FOR EACH ROW
```

```
-- FOR EACH ROW means that it is a row trigger
```

Trigger type: row

Action

```
EXECUTE PROCEDURE CheckBudget ();
```

Trigger condition

Create or Replace Trigger Statement

Implementation of action

```
CREATE OR REPLACE FUNCTION CheckBudget ()  
RETURNS TRIGGER AS  
$$  
BEGIN
```

Start of trigger action

Variables `OLD` and `NEW` represent a row before modification or deletion and a row after modification or insertion

```
IF NOT ( NEW.budget BETWEEN 1 AND 7000 ) THEN
```

Application of NEW variable in a row trigger

Abnormal termination of a trigger together with a transaction that fired a trigger

```
RAISE EXCEPTION 'Budget of department % cannot be equal to %',  
                NEW.name, NEW.budget;  
END IF;
```

Abnormal termination of a trigger

Create or Replace Trigger Statement

Normal termination of a trigger together with a transaction that fired a trigger

```
RETURN NEW;
```

Normal termination of a trigger

End of trigger body

```
END;
```

```
$$
```

```
LANGUAGE plpgsql;
```

End of trigger action

Create or Replace Trigger Statement

A complete `CREATE OR REPLACE TRIGGER` statement

```
CREATE OR REPLACE TRIGGER CheckBudget
BEFORE UPDATE OF budget ON DEPARTMENT
FOR EACH ROW
EXECUTE PROCEDURE CheckBudget ();
```

A sample row trigger

Implementation of an action

```
CREATE OR REPLACE FUNCTION CheckBudget ()
RETURNS TRIGGER AS
$$
BEGIN
    IF NOT ( NEW.budget BETWEEN 1 AND 7000 ) THEN
        RAISE EXCEPTION 'Budget of department % cannot be equal to %',
            NEW.name, NEW.budget;
    END IF;
    RETURN NEW;
END;
$$
LANGUAGE plpgsql;
```

Action

Create or Replace Trigger Statement

Another implementation of the same trigger with **WHEN** clause

```
CREATE OR REPLACE TRIGGER CheckBudget
  BEFORE UPDATE OF budget ON DEPARTMENT
  FOR EACH ROW
  WHEN ( NOT ( NEW.budget BETWEEN 1 AND 7000 ) )
  EXECUTE PROCEDURE AbortUpdate();
```

A sample row trigger

Implementation of an action

```
CREATE OR REPLACE FUNCTION AbortUpdate()
  RETURNS TRIGGER AS
  $$
  BEGIN
    RAISE EXCEPTION 'Budget of department % cannot be equal to %',
      NEW.name, NEW.budget;
    RETURN NEW;
  END;
  $$
LANGUAGE plpgsql;
```

Action

Create or Replace Trigger Statement

The results from processing when a trigger detects an incorrect value

```
UPDATE DEPARTMENT
SET budget = 100000
WHERE name = 'Physics';
ERROR:  Budget of department Physics cannot be equal to 100000.0
CONTEXT:  PL/pgSQL function abortupdate() line 3 at RAISE;
```

A trigger detects an incorrect value

The results from processing when a trigger does not detect an incorrect value

```
UPDATE DEPARTMENT
SET budget = 100
WHERE name = 'Physics';
UPDATE 1
```

A trigger does not detect an incorrect value

Create or Replace Trigger Statement

The following **temporal options** are available

- **BEFORE** - a trigger fires before a triggering event
- **AFTER** - a trigger fires after a triggering event
- **INSTEAD OF** - a trigger fires instead of a triggering event, it is typically used to correctly implement **view update** operation i.e. a correct modification of **base relational tables** through an update performed on a **relational view**

Sample applications of **temporal options**

Fire a trigger before **UPDATE** operation on a relational table **DEPARTMENT**

```
BEFORE UPDATE ON DEPARTMENT
```

A sample temporal option

Fire a trigger after **DELETE** operation on a relational table **COURSE**

```
AFTER DELETE ON COURSE
```

A sample temporal option

Create or Replace Trigger Statement

Fire a trigger instead of `INSERT` operation on a relational view `EMPVIEW`

```
INSTEAD OF INSERT ON EMPVIEW
```

A sample temporal option

Fire a trigger after `INSERT` into a table `DEPARTMENT` or after `DELETE` from a table `DEPARTMENT`

```
AFTER INSERT OR DELETE ON DEPARTMENT
```

A sample temporal option

Fire a trigger before `UPDATE` operation on a column `budget` in a relational table `DEPARTMENT`

```
BEFORE UPDATE OF budget ON DEPARTMENT
```

A sample temporal option

Create or Replace Trigger Statement

The following events can fire a trigger

- INSERT
- DELETE
- UPDATE
- TRUNCATE

Sample applications of **DML events**

```
BEFORE UPDATE OF cnum, title ON COURSE
```

A sample DML event

```
AFTER INSERT ON DEPARTMENT
```

A sample DML event

```
INSTEAD OF TRUNCATE ON DEPARTMENT
```

A sample DML event

```
AFTER DELETE OR INSERT OR UPDATE ON COURSE
```

A sample DML event

Create or Replace Trigger Statement

Condition determines whether a trigger processes its body after it has been fired

Sample applications of **condition**

```
WHEN (NEW.budget > 7000)
```

A sample condition

```
WHEN (OLD.budget < 1000 AND NEW.budget > 2000);
```

A sample condition

```
WHEN (NEW.amount > 1000 );
```

A sample condition

```
WHEN (OLD.credits IN (6, 12));
```

A sample condition

OLD and **NEW** are **record variable** such that for

- **INSERT** triggering operation **OLD** contains no values and **NEW** contains the new values
- **UPDATE** triggering operation **OLD** contains the old values and **NEW** contains the new values
- **DELETE** triggering operation **OLD** contains the old values and **NEW** contains no values

Stored PL/pgSQL

Statement
Database
Triggers



Statement Database Trigger

A **statement trigger** fires once either before or after a triggering event

A sample **statement** trigger records in a relational table **DELETE** and **INSERT** operations on a relational table

Trigger name	
CREATE OR REPLACE TRIGGER Audit	
Temporal option and event specification	
AFTER DELETE OR INSERT ON DEPARTMENT	
FOR EACH STATEMENT	
Action	
EXECUTE FUNCTION InsertAudit();	

A function **InsertAudit** inserts into a relational table **AUDIT** information about **DELETE** and **INSERT** operations together with the timestamps

Table AUDIT	
CREATE TABLE AUDIT (	
statement VARCHAR(10) NOT NULL,	
date_time TIMESTAMP NOT NULL,	
CONSTRAINT AUDIT_PK PRIMARY KEY (date_time));	

Statement Database Trigger

A variable `TG_OP` contains a name of operation that fired a trigger

```
CREATE OR REPLACE FUNCTION InsertAudit()  
RETURNS TRIGGER AS  
$$  
BEGIN  
    INSERT INTO AUDIT VALUES (TG_OP, current_timestamp);  
    RETURN NULL;  
END;  
$$  
LANGUAGE plpgsql;
```

A function InsertAudit

Other special variables:

- `NEW`, `OLD`,
- `TG_NAME`: a name of trigger
- `TG_WHEN`: either `BEFORE` or `AFTER` or `INSTEAD`
- `TG_LEVEL`: either `ROW` or `STATEMENT` trigger

Statement Database Trigger

Testing with a single insertion, single deletion and many deletions

Testing

```
INSERT INTO DEPARTMENT VALUES ('Quantum Computers','QC',30,'Kate',223426.9 );  
DELETE FROM DEPARTMENT WHERE name = 'Arts';  
DELETE FROM DEPARTMENT;
```

Rows inserted into AUDIT table

```
SELECT * FROM AUDIT;
```

statement	date_time
INSERT	2023-12-06 14:08:20.475765
DELETE	2023-12-06 14:08:33.19901
DELETE	2023-12-06 14:08:41.729328



Stored PL/pgSQL Row Database Trigger

Row Database Trigger

A **row trigger** fires either after or before a triggering event affects a row in a relational table

- When a **temporal option** **BEFORE** is used a trigger fires **once before** a triggering event affects a row in a relational table
- When a **temporal option** **AFTER** is used a trigger fires **once after** a triggering event affects a row in a relational table

For example, if a **temporal option** and **event** are

BEFORE INSERT ON DEPARTMENT

A temporal option and event

then a trigger fires before each insertion into a relational table (it is possible to have many insertions when a multirow **INSERT** statement is processed)

For example, if a **temporal option** and **event** are

AFTER UPDATE ON EMPLOYEE

A temporal option and event

then a trigger fires after a row is updated in a relational table, if a triggering event updates **n** rows then a trigger fires **n** times

Row Database Trigger

For example, if a **temporal option** and **event** are

```
AFTER DELETE ON PROJECT
```

A temporal option and event

Then a trigger fires after a row is deleted from a relational table, if a triggering event deletes **n** rows then a trigger fires **n** times

Row Database Trigger

A sample **row** trigger that records in a relational table **AUDIT** information about each row updated in a relational table **DEPARTMENT**

```
CREATE OR REPLACE TRIGGER UpdateDepartment
```

Trigger name

```
AFTER UPDATE ON DEPARTMENT
```

Temporal option and event specification

```
FOR EACH ROW
```

Action

```
EXECUTE FUNCTION UpdateDepartment ();
```

A function **UpdateDepartment** inserts into a relational table **AUDIT** information about **UPDATE** operation, timestamp, name of a department, old and new value of budget

```
CREATE TABLE AUDIT (  
  statement VARCHAR(10) NOT NULL,  
  date_time TIMESTAMP NOT NULL,  
  name VARCHAR(50) NOT NULL,  
  oldbudget DECIMAL(9,1) NOT NULL,  
  newbudget DECIMAL(9,1) NOT NULL);
```

Table AUDIT

Row Database Trigger

A variable `TG_OP` contains a name of operation that fired a trigger

A variable `OLD.name` contains a name of departemnt

A variable `OLD.budget` contains a value of budget befor update and a variable `NEW.budget` contains a value of budget after update

A function UpdateDepartment

```
CREATE OR REPLACE FUNCTION UpdateDepartment ()
RETURNS TRIGGER AS
$$
BEGIN
    INSERT INTO AUDIT VALUES (TG_OP, current_timestamp, OLD.name, OLD.budget,
                                NEW.budget);

    RETURN NULL;
END;
$$
LANGUAGE plpgsql;
```

Row Database Trigger

Testing with a single update of many rows in a relational table

DEPARTMENT

Testing

```
UPDATE DEPARTMENT SET budget = budget + 1000;  
UPDATE 7
```

Rows inserted into AUDIT table

```
SELECT * FROM AUDIT;
```

statement	date_time	name	oldbudget	newbudget
UPDATE	2023-12-06 18:56:03.402062	Computer Science	223426.9	224426.9
UPDATE	2023-12-06 18:56:03.402062	Physics	223453.9	224453.9
UPDATE	2023-12-06 18:56:03.402062	Engineering	313455.9	314455.9
UPDATE	2023-12-06 18:56:03.402062	Mathematics	133456.9	134456.9
UPDATE	2023-12-06 18:56:03.402062	Education	173557.9	174557.9
UPDATE	2023-12-06 18:56:03.402062	Science	163459.9	164459.9
UPDATE	2023-12-06 18:56:03.402062	Quantum Computers	223426.9	224426.9

(7 rows)



Database Trigger

Sample Application

Sample Application

We consider a database created through processing of the following `CREATE TABLE` statements

Relational table DEPARTMENT

```
CREATE TABLE DEPARTMENT (  
  name          VARCHAR(50)          NOT NULL,  
  code          CHAR(5)              NOT NULL,  
  total_staff_number DECIMAL(2)      NOT NULL,  
  chair         VARCHAR(50)          NULL,  
  budget        DECIMAL(9,1)         NOT NULL,  
  CONSTRAINT dept_pkey PRIMARY KEY(name),  
  CONSTRAINT dept_cke1 UNIQUE(code),  
  CONSTRAINT dept_cke2 UNIQUE(chair),  
  CONSTRAINT dept_check1 CHECK (total_staff_number BETWEEN 1 AND 50) );
```

Relational table COURSE

```
CREATE TABLE COURSE (  
  cnum          CHAR(7)              NOT NULL,  
  title         VARCHAR(200)         NOT NULL,  
  credits       DECIMAL(2)           NOT NULL,  
  offered_by    VARCHAR(50)          NULL,  
  CONSTRAINT course_pkey PRIMARY KEY(cnum),  
  CONSTRAINT course_check1 CHECK (credits IN (6, 12)),  
  CONSTRAINT course_fkey1 FOREIGN KEY(offered_by)  
    REFERENCES DEPARTMENT(name) ON DELETE CASCADE );
```

Sample Application

We would like to add information about the total number of courses taught by each department

A correct modification of a design adds a column `total_courses` to a relational table `DEPARTMENT`

Addition of a column `total_courses` does not violate BCNF for a relational table `DEPARTMENT`

There is no need to add any new relational table to the design

Adding a column `total_courses` to a table `DEPARTMENT`

```
ALTER TABLE DEPARTMENT ADD total_courses DECIMAL(2);
```

A column `total_courses` is filled with data in the following way

Filling a column `total_courses` with data

```
UPDATE DEPARTMENT
SET total_courses = (SELECT COUNT(*)
                    FROM COURSE
                    WHERE course.offered_by = DEPARTMENT.name);
```


Sample Application

Each time a course is added or deleted or moved to another department a statement trigger supposed to update the values in a column `total_courses` in a relational table `DEPARTMENT`

A statement trigger that updates a column `total_courses` in a table `DEPARTMENT`

```
CREATE OR REPLACE TRIGGER UpdateTotal
AFTER DELETE OR INSERT OR UPDATE ON COURSE
FOR EACH STATEMENT
EXECUTE FUNCTION ChangeTotal();
```

A function `ChangeTotal()` is implemented in the following way

Updating the values in a column `total_courses` in a relational table `DEPARTMENT`

```
CREATE OR REPLACE FUNCTION ChangeTotal()
RETURNS TRIGGER AS
$$
BEGIN
    UPDATE DEPARTMENT
    SET total_courses = (SELECT COUNT(*)
                        FROM COURSE
                        WHERE course.offered_by = DEPARTMENT.name);

    RETURN NULL;
END;
$$ LANGUAGE plpgsql;
```

Sample Application

The same task can be implemented as a row trigger in the following way

A row trigger that updates a column `total_courses` in a table `DEPARTMENT`

```
CREATE OR REPLACE TRIGGER UpdateTotal
AFTER DELETE OR INSERT OR UPDATE OF offered_by ON COURSE
FOR EACH ROW
EXECUTE FUNCTION RowChangeTotal();
```

A function `RowChangeTotal()` is implemented in the following way

Updating the values in a column `total_courses` in a relational table `DEPARTMENT`

```
CREATE OR REPLACE FUNCTION RowChangeTotal()
RETURNS TRIGGER AS $$
BEGIN
    IF (TG_OP = 'DELETE') THEN
        UPDATE DEPARTMENT SET total_courses = total_courses - 1 WHERE name = OLD.offered_by;
    ELSIF (TG_OP = 'INSERT') THEN
        UPDATE DEPARTMENT SET total_courses = total_courses + 1 WHERE name = NEW.offered_by;
    ELSE
        UPDATE DEPARTMENT SET total_courses = total_courses - 1 WHERE name = OLD.offered_by;
        UPDATE DEPARTMENT SET total_courses = total_courses + 1 WHERE name = NEW.offered_by;
    END IF;
    RETURN NULL;
END;
$$ LANGUAGE plpgsql;
```

Database Trigger

Problems With Database
Triggers



Problems With Database Triggers

Infinite chains of trigger invocations

- What to do when a trigger **A** while processing its body fires a trigger **B** and a trigger **B** while processing its body fires a trigger **A** ?

Indeterministic trigger invocations

- It may happen that due to a database transaction serialization mechanisms the same chain of trigger invocations will be processed (serialized) in many different way by a transaction scheduler, e.g. if two triggers **A** and **B** fire in more or less the same moment in time then sometimes **A** will be processed before **B** and sometimes **B** will be processed before **A**

Lack of external control

- Long chains of trigger invocations contribute to very serious data security risks, e.g. it is possible to "hide" malicious code at the end of long chains of trigger invocations

Lack of design methodology

- The ad hoc uncontrolled and not well planned additions of new triggers lead to a situation where after addition or modification of a trigger there is no certainty that the chains of trigger invocations do not corrupt a database

References

[PostgreSQL 16.1 Documentation, Chapter 39. Triggers](#)

[PostgreSQL 16.1 Documentation, CREATE TRIGGER](#)

[PostgreSQL Tutorial, PostgreSQL Triggers](#)

T. Connolly, C. Begg, Database Systems, A Practical Approach to Design, Implementation, and Management, Chapter 8 Advanced SQL, Pearson Education Ltd, 2015